CSE 130 - Fall 2020 - Section 01 - Assignment 1

Due: Monday, November 2 at 9:00PM

Goals

The goal for Assignment 1 is to implement a simple single-threaded HTTP server. The server will respond to simple GET and PUT commands to read and write "files" named by 10-character ASCII names. The server will persistently store files in a directory on the server, so it can be restarted or otherwise run on a directory that already has files. As usual, you must have a design document along with your `README.md` in your `git`. Your code must build `httpserver` using `make`.

Programming assignment: HTTP server

Design document

Before writing code for this assignment, as with every other assignment, you must write up a design document. Your design document must be called `DESIGN.pdf`, and must be in PDF (you can easily convert other document formats, including plain text, to PDF).

Your design should describe the design of your code in enough detail that a knowledgeable programmer could duplicate your work. This includes descriptions of the data structures you use, non-trivial algorithms and formulas, and a description of each function with its purpose, inputs, outputs, and assumptions it makes about inputs or outputs.

Write your design document *before* you start writing code. It'll make writing code a lot easier. Also, if you want help with your code, the first thing we're going to ask for is your design document. We're happy to help you with the design, but we can't debug code without a design any more than you can.

TESTING AND ASSIGNMENT QUESTION

In the design document, you will also describe the testing you did on your program and answer any short questions below. The testing can be unit testing (testing of individual functions or smaller pieces of the program) or whole-system testing, which involves running your code in particular scenarios.

For Assignment 1, please answer the following question:

- What happens in your implementation if, during a PUT with a `Content-Length`, the connection was closed, ending the communication early? This extra concern was not present in your implementation of `dog`. Why not? Hint: this is an example of complexity being added by an extension of requirements (in this case, data transfer over a network).

Program functionality

You may not use standard libraries for HTTP; you have to implement this yourself. You may use standard networking (and file system) system calls, but not any `FILE *` calls except for printing

to the screen (*e.g.,* error messages). Note that string functions like `sprintf()` and `sscanf()` aren't `FILE *` calls.

Your code may be either C or C++, but all source files must have a `.cpp` suffix and be compiled with no errors or warnings using the following flags: `-std=gnu++11 -Wall -Wextra -Wpedantic -Wshadow`

HTTP protocol

The HTTP protocol is used by clients and servers for a significant amount of web communication. It's designed to be simple, easy to parse (in code), and easy to read (by a human). Your server will need to send and receive files via `http`, so the best approach may be to reuse your code for copying data between file descriptors (which you should be well-acquainted with after `asgn0`!).

Your server will need to be able to parse simple HTTP headers; you're encouraged to use string functions (but not `FILE *` functions) to do this. The `http` protocol that you need to implement is very simple. The client sends a request to the server that either asks to send a file from client to server (PUT) or fetch a file from server to client (GET).

An HTTP header consists of one or more lines of (ASCII) text, followed by a blank (empty) line. The first line of the header is a *single line* specifying the action. A PUT header looks like this (note the blank line at the end):

```
PUT /0123456789 HTTP/1.1\r\nContent-Length: 460\r\n\r\n
```

The newlines are encoded as `\r\n`, and the data being sent immediately follows the header. The sent data may include any bytes of data, including NUL (`\0`). The `Content-Length` header line is optional; the client may send it to indicate how much data follows the header. if it's included, the server should stop reading data after the specified number of bytes, and look for another header (if the client hasn't closed the connection). If there's no `Content-Length` header, the server copies data until the client closes the connection (`read()` reads end-of-file). For the above example, the name that the server will bind to the data is `0123456789`.

For GET, the header looks like this. There's no `Content-Length` header because the client doesn't know the length of the content (again, notice the blank line at the end):

```
GET /ABCDEFabcd HTTP/1.1\r\n\r\n
```

All valid resource names in this assignment *must* be 10 ASCII characters long, and must consist *only* of the upper and lower case letters of the (English) alphabet (52 characters), and the digits 0–9 (10 characters). If a request includes an invalid name, the server must fail the request and respond accordingly.

The server must respond to a PUT or GET with a "response", which is a response header optionally followed by data. An example response header looks like this:

```
HTTP/1.1 200 OK\r\nContent-Length: 0\r\n\r\n
```

The `200` is a status code—200 means "OK". The `OK` message is an informational description of the code. For example, the `404` status code could say "File not found". The server must fill in the appropriate status code and message. For a response to a GET request (and other requests), the server must provide a `Content-Length:` line in the header, similar to the one shown in the PUT request. The header is followed by a blank line and, for a GET response, by the data that the client has requested. As before, the data may include *any* data, including `NUL` bytes. (A response to a put request would have a content length of 0.)

You can find a list of HTTP status codes at:

https://en.wikipedia.org/wiki/List_of_HTTP_status_codes.

The only status codes you'll *need* to implement are 200 (OK), 201 (Created), 400 (Bad Request), 403 (Forbidden), 404 (Not Found), and 500 (Internal Server Error). You may use additional status codes if you like, but these are the only required ones. Look at the link above to determine when to use each one.

Your server will need to be able to handle malformed and erroneous requests and respond appropriately, without crashing. Note that a "bad" name is *not* the same thing as a valid name that doesn't correspond to an existing file. You may assume that a header will be no longer than 4 KiB, though the data that follows it (for a PUT) may be much longer. Similarly, response headers will be less than 4 KiB, but the data may be (much) longer.

HTTP server

Your server binary must be called `httpserver`. Your HTTP server is a single-threaded server that will listen on a user-specified port and respond to HTTP PUT and GET requests on that port. The address to listen to and the port number are specified on the command line. (Yes, it *is* necessary to specify the server address, since your computer has multiple Internet addresses, including `localhost`.)

The first argument to `httpserver` is the address of the HTTP server to contact, which may be specified as a hostname or an IP address; your software must handle either one. The second, optional, argument to `httpserver` is the port number on which to listen. If there's no second argument, assume the standard HTTP port, port 80.

Your server will use the directory in which it's run to `write(2)` files that are PUT, and `read(2)` files for which a GET request is made. **All file I/O for user data must be done via `read()` and `write()`.**

Testing your code

You should test your code on your own system. You can run the server on `localhost` using a port number above 1024 (*e.g.*, 8888). Come up with requests you can make of your server, and try them using `curl(1)`.

You might also consider cloning a new copy of your repository (from GitLab@UCSC) to a clean directory to see if it builds properly, and runs as you expect. That's an easy way to tell if your repository has all of the right files in it. You can then delete the newly-cloned copy of the directory on your local machine once you're done with it.

README

As for previous assignments, your repository must include (`README.md`). The `README.md` file should be short, and contain any instructions necessary for running your code. You should also

list limitations or issues in `README.md`, telling a user if there are any known issues with your code. You should also write the team members in the README file.

Submitting your assignment

All of your files for Assignment 1 must be in the `asgn1` directory in your `git`. When you push your repository to GitLab@UCSC, make sure to include the following:

- There are no "bad" files in the `asgn1` (*i.e.,* object files).

- Your assignment builds in `asgn1 make` to produce `httpserver`.

- All required files (source files, `DESIGN.pdf`, `README.md`) are present in `asgn1`. Note that you do not have to write a client program since you can use curl as a client.

Hints

- Start early on the design. This is a more difficult program than `dog`!

- You'll need to use (at least) the system calls `socket`, `bind`, `listen`, `accept`, `connect`, `send`, `recv`, `open`, `read`, `write`, `close`. The last four calls should be familiar from Assignment 0, and `send` and `recv` are very similar to `write` and `read`, respectively. You might also want to investigate `dprintf(3)` for printing to a file descriptor and `sscanf(3)` for parsing data in a string (*i.e.,* a buffer). You should read the `man(ual)` pages or other documentation for these functions. Don't worry about the complexity of opening a socket; we'll discuss it in section. **You may not use any calls for operating on files or network sockets other than those above.**

- Test your server using an existing Web client (we recommend `curl(1)`). Make sure you test error conditions as well as "normal" operation.

- Aggressively check for and report errors via a response. If your server runs into a problem well into sending data in response to a `GET`, you may not be able to send an error header (the header may have been sent long ago). Instead, you should just close the connection. Normally, however, responses are your server's way of notifying the client of an error.

* When using curl, add a / in front of the file name in the request, to make sure that curl performs the default behavior

* Start with these curl commands for testing:

    * curl -T t1 http://localhost:8080/0123456789 -v

    * curl http://localhost:8080/0123456789 -v

- Your commit must contain the following files:

    - `README.md`

- DESIGN.pdf

- Makefile

- source file(s) for the server (see update above)

It may *not* contain any .o files. You may, if you wish, include the "source" files for your DESIGN.pdf in your repo, but you don't have to. After running make, your directory must contain httpserver. Your source files must be .cpp files (and the corresponding headers, if needed).

- If you need help, use online documentation such as man pages and documentation on Makefiles. If you still need help, ask the course staff.

Grading

As with all of the assignments in this class, we will be grading you on *all* of the material you turn in, with the *approximate* distribution of points as follows: design document and answer to assignment question (30%); coding practices (10%); functionality (60%).

Your code must compile to be graded. A submission that cannot compile may receive a maximum grade of as low as 5%.